

# Project 1 - BBC News Classification

In [49]:

```
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import nltk
import re
from nltk.corpus import stopwords
from nltk.tokenize import word_tokenize
from wordcloud import WordCloud, STOPWORDS
from sklearn.model_selection import train_test_split
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, classification_report
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.naive_bayes import GaussianNB
```

## Importing Data

In [2]:

```
data = pd.read_csv("D:\Python Data Science\BBC News Classification\BBC News Train.csv")
data
```

Out[2]:

	ArticleId	Text	Category
0	1833	worldcom ex-boss launches defence lawyers defe...	business
1	154	german business confidence slides german busin...	business
2	1101	bbc poll indicates economic gloom citizens in ...	business
3	1976	lifestyle governs mobile choice faster bett...	tech
4	917	enron bosses in \$168m payout eighteen former e...	business
...	...	...	...
1485	857	double eviction from big brother model caprice...	entertainment
1486	325	dj double act revamp chart show dj duo jk and ...	entertainment
1487	1590	weak dollar hits reuters revenues at media gro...	business
1488	1587	apple ipod family expands market apple has exp...	tech
1489	538	santy worm makes unwelcome visit thousands of ...	tech

1490 rows × 3 columns

## Understanding Features & Target Attributes

In [3]:

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1490 entries, 0 to 1489
Data columns (total 3 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   ArticleId   1490 non-null   int64
 1   Text        1490 non-null   object
 2   Category    1490 non-null   object
dtypes: int64(1), object(2)
memory usage: 35.0+ KB
```

In [4]:

```
data['Category'].value_counts()
```

Out[4]:

```
sport          346
business       336
politics       274
entertainment  273
tech           261
Name: Category, dtype: int64
```

## Checking Any Null Values

In [5]:

```
data.isnull().any()
```

Out[5]:

```
ArticleId    False
Text         False
Category     False
dtype: bool
```

## Countplot of Target Attribute(Category)

In [6]:

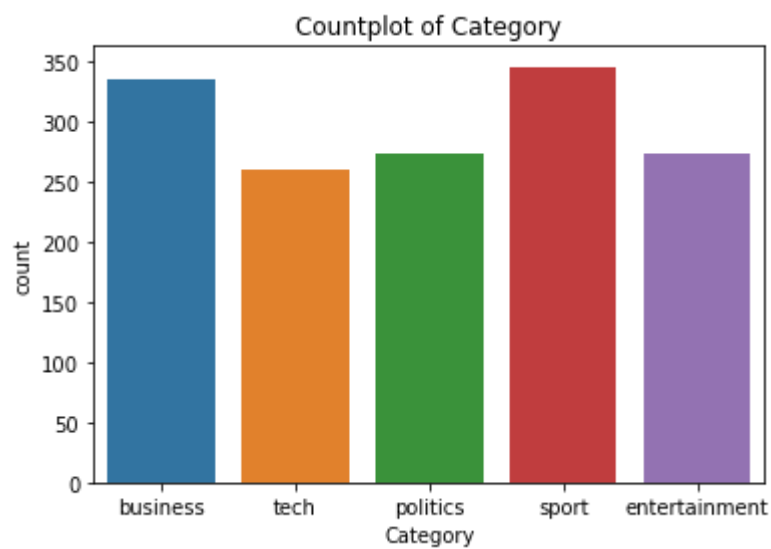
```
sns.countplot(data.Category).set_title('Countplot of Category')
```

C:\Users\Shivam\anaconda3\lib\site-packages\seaborn\\_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
```

Out[6]:

```
Text(0.5, 1.0, 'Countplot of Category')
```



In [8]:

```
data['Text_length'] = data['Text'].str.len()
data
```

Out[8]:

	ArticleId	Text	Category	Text_length
0	1833	worldcom ex-boss launches defence lawyers defe...	business	1866
1	154	german business confidence slides german busin...	business	2016
2	1101	bbc poll indicates economic gloom citizens in ...	business	3104
3	1976	lifestyle governs mobile choice faster bett...	tech	3618
4	917	enron bosses in \$168m payout eighteen former e...	business	2190
...	...	...	...	...
1485	857	double eviction from big brother model caprice...	entertainment	1266
1486	325	dj double act revamp chart show dj duo jk and ...	entertainment	3111
1487	1590	weak dollar hits reuters revenues at media gro...	business	1370
1488	1587	apple ipod family expands market apple has exp...	tech	3242
1489	538	santy worm makes unwelcome visit thousands of ...	tech	1723

1490 rows × 4 columns

# Distribution Plot

In [9]:


```
sns.distplot(data.Text_length).set_title('News length Distribution')
```

C:\Users\Shivam\anaconda3\lib\site-packages\seaborn\distributions.py:2551: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

warnings.warn(msg, FutureWarning)

Out[9]:

Text(0.5, 1.0, 'News length Distribution')



## Constructing Wordcloud

In [10]:

```
def text_wordcloud(words):  
    wc = WordCloud(width=800, height=500, max_font_size=110).generate(words)  
    plt.figure(figsize=(10, 7))  
    plt.axis('off')  
    plt.imshow(wc, interpolation="bilinear")
```

```
temp_sport = data[data['Category'] == "sport"]
sport_sentences = temp_sport['Text'].tolist()
text = " ".join(sport_sentences)

text_wordcloud(text)
```



```
temp_sport = data[data['Category'] == "entertainment"]
entertainment_sentences = temp_sport['Text'].tolist()
text = " ".join(entertainment_sentences)

text_wordcloud(text)
```



```
temp_sport = data[data['Category'] == "politics"]
politics_sentences = temp_sport['Text'].tolist()
text = " ".join(politics_sentences)

text_wordcloud(text)
```





```
temp_sport = data[data['Category'] == "business"]
business_sentences = temp_sport['Text'].tolist()
text = " ".join(business_sentences)

text_wordcloud(text)
```



```
temp_sport = data[data['Category'] == "tech"]
tech_sentences = temp_sport['Text'].tolist()
text = " ".join(tech_sentences)

text_wordcloud(text)
```





# Data Preprocessing

Removing special characters

- \r
- \n

Removing Punctuations & Stopwords

In [16]:

```
def preprocess_text(text):
    text = text.lower().replace('\n', ' ').replace('\r', '').strip()
    text = re.sub(' +', ' ', text)
    text = re.sub(r'^\w\s', '', text)

    mystopwords = set(stopwords.words('english'))
    word_tokenize_output = word_tokenize(text)

    temp_word_list = []
    for each_word in word_tokenize_output:
        if each_word not in mystopwords:
            temp_word_list.append(each_word)
    temp_word_list = " ".join(temp_word_list)

    return temp_word_list
```

In [17]:

```
data['Cleaned_text'] = data['Text'].apply(preprocess_text)
data['Cleaned_text_length'] = data['Cleaned_text'].str.len()
data
```

Out[17]:

	ArticleId	Text	Category	Text_length	Cleaned_text	Cleaned_text_length
0	1833	worldcom ex-boss launches defence lawyers defe...	business	1866	worldcom exboss launches defence lawyers defen...	1423
1	154	german business confidence slides german busin...	business	2016	german business confidence slides german busin...	1508
2	1101	bbc poll indicates economic gloom citizens in ...	business	3104	bbc poll indicates economic gloom citizens maj...	2145
3	1976	lifestyle governs mobile choice faster bett...	tech	3618	lifestyle governs mobile choice faster better ...	2412
4	917	enron bosses in \$168m payout eighteen former e...	business	2190	enron bosses 168m payout eighteen former enron...	1581
...	...	...	...	...	...	...
1485	857	double eviction from big brother model caprice...	entertainment	1266	double eviction big brother model caprice holb...	869
1486	325	dj double act revamp chart show dj duo jk and ...	entertainment	3111	dj double act revamp chart show dj duo jk joel...	2080
1487	1590	weak dollar hits reuters revenues at media gro...	business	1370	weak dollar hits reuters revenues media group ...	990
1488	1587	apple ipod family expands market apple has exp...	tech	3242	apple ipod family expands market apple expande...	2273
1489	538	santy worm makes unwelcome visit thousands of ...	tech	1723	santy worm makes unwelcome visit thousands web...	1225

1490 rows × 6 columns

## Label Encoding

In [18]:

```
from sklearn import preprocessing
label_encoder = preprocessing.LabelEncoder()
data['Category_target'] = label_encoder.fit_transform(data['Category'])
```

In [19]:

data

Out[19]:

	ArticleId	Text	Category	Text_length	Cleaned_text	Cleaned_text_length	Category
0	1833	worldcom ex-boss launches defence lawyers defe...	business	1866	worldcom exboss launches defence lawyers defen...	1423	
1	154	german business confidence slides german busin...	business	2016	german business confidence slides german busin...	1508	
2	1101	bbc poll indicates economic gloom citizens in ...	business	3104	bbc poll indicates economic gloom citizens maj...	2145	
3	1976	lifestyle governs mobile choice faster bett...	tech	3618	lifestyle governs mobile choice faster better ...	2412	
4	917	enron bosses in \$168m payout eighteen former e...	business	2190	enron bosses 168m payout eighteen former enron...	1581	
...	...	...	...	...	...	...	
1485	857	double eviction from big brother model caprice...	entertainment	1266	double eviction big brother model caprice holb...	869	
1486	325	dj double act revamp chart show dj duo jk and ...	entertainment	3111	dj double act revamp chart show dj duo jk joel...	2080	
1487	1590	weak dollar hits reuters revenues at media gro...	business	1370	weak dollar hits reuters revenues media group ...	990	
1488	1587	apple ipod family expands market apple has exp...	tech	3242	apple ipod family expands market apple expande...	2273	

ArticleId	Text	Category	Text_length	Cleaned_text	Cleaned_text_length	Category
1489	538 santy worm makes unwelcome visit thousands of ...	tech	1723	santy worm makes unwelcome visit thousands web...	1225	

1490 rows × 7 columns

## Splitting Data in Training Set and Testing Set

In [20]:

```
x_train,x_test,y_train,y_test = train_test_split(data['Cleaned_text'], data['Category_target'],
```

In [21]:

```
print(x_train.shape)
print(y_train.shape)
print(x_test.shape)
print(y_test.shape)
```

(1117,)

(1117,)

(373,)

(373,)

In [28]:

```
tfidf = TfidfVectorizer(encoding = 'utf-8',
                        ngram_range = (1,2),
                        max_df = 1.,
                        min_df = 10,
                        max_features = 300,
                        )

features_train = tfidf.fit_transform(x_train).toarray()
labels_train = y_train
print(features_train.shape)

features_test = tfidf.transform(x_test).toarray()
labels_test = y_test
print(features_test.shape)
```

(1117, 300)

(373, 300)

## Building a Model

## Logistic Regression

In [34]:

```
lc_model = LogisticRegression()
lc_model.fit(features_train, labels_train)
lc_model_predictions = lc_model.predict(features_test)

print("Accuracy ->", accuracy_score(labels_test, lc_model_predictions))
print(classification_report(labels_test, lc_model_predictions))
```

Accuracy -&gt; 0.9544235924932976

	precision	recall	f1-score	support
0	0.95	0.94	0.95	86
1	0.97	0.97	0.97	62
2	0.96	0.96	0.96	75
3	0.97	0.97	0.97	79
4	0.92	0.93	0.92	71
accuracy			0.95	373
macro avg	0.95	0.95	0.95	373
weighted avg	0.95	0.95	0.95	373

## Decision Tree

In [38]:

```
dc_model = DecisionTreeClassifier()
dc_model.fit(features_train, labels_train)
dc_model_predictions = dc_model.predict(features_test)

print("Accuracy ->", accuracy_score(labels_test, dc_model_predictions))
print(classification_report(labels_test, dc_model_predictions))
```

Accuracy -&gt; 0.7855227882037533

	precision	recall	f1-score	support
0	0.71	0.76	0.73	86
1	0.85	0.74	0.79	62
2	0.74	0.77	0.76	75
3	0.78	0.92	0.84	79
4	0.91	0.72	0.80	71
accuracy			0.79	373
macro avg	0.80	0.78	0.79	373
weighted avg	0.79	0.79	0.79	373

## Random Forest

In [40]:

```
rf_model = RandomForestClassifier()
rf_model.fit(features_train, labels_train)
rf_model_predictions = rf_model.predict(features_test)

print("Accuracy ->", accuracy_score(labels_test, rf_model_predictions))
print(classification_report(labels_test, rf_model_predictions))
```

Accuracy -&gt; 0.9302949061662198

	precision	recall	f1-score	support
0	0.95	0.91	0.93	86
1	0.90	0.90	0.90	62
2	0.93	0.95	0.94	75
3	0.92	0.97	0.94	79
4	0.94	0.92	0.93	71
accuracy			0.93	373
macro avg	0.93	0.93	0.93	373
weighted avg	0.93	0.93	0.93	373

## KNeighborsClassifier

In [44]:

```
kn_model = KNeighborsClassifier()
kn_model.fit(features_train, labels_train)
kn_model_predictions = kn_model.predict(features_test)

print("Accuracy ->", accuracy_score(labels_test, kn_model_predictions))
print(classification_report(labels_test, kn_model_predictions))
```

Accuracy -&gt; 0.900804289544236

	precision	recall	f1-score	support
0	0.87	0.87	0.87	86
1	0.95	0.84	0.89	62
2	0.84	0.92	0.88	75
3	0.97	0.95	0.96	79
4	0.89	0.92	0.90	71
accuracy			0.90	373
macro avg	0.90	0.90	0.90	373
weighted avg	0.90	0.90	0.90	373

## Naive Bayes



In [47]:

```
nb_model = GaussianNB()
nb_model.fit(features_train, labels_train)
nb_model_predictions = nb_model.predict(features_test)

print("Accuracy ->", accuracy_score(labels_test, nb_model_predictions))
print(classification_report(labels_test, nb_model_predictions))
```

Accuracy -> 0.9088471849865952

	precision	recall	f1-score	support
0	0.90	0.86	0.88	86
1	0.87	0.87	0.87	62
2	0.92	0.93	0.93	75
3	0.96	0.97	0.97	79
4	0.88	0.90	0.89	71
accuracy			0.91	373
macro avg	0.91	0.91	0.91	373
weighted avg	0.91	0.91	0.91	373

In [ ]: